

CSCI 509 - Operating Systems, Fall 2019

Assignment 2: Threads and Synchronization

Due Date: Friday, Oct 11, 2019

Points: 200

1 Overview and Goal

In this assignment you will review threads and gain more experience writing programs involving concurrency control. You will begin by studying the thread package, which implements multi-threading. You will make modifications and additions to the existing code. Then you will use the threads package to solve several concurrency problems.

In addition, you will gain familiarity programming in the KPL language.

If you have difficulties with this assignment or want to go into more depth, take a look at the document called "The Thread Scheduler and Concurrency Control Primitives":

<http://facultyweb.cs.wvu.edu/~phil/classes/blitz/BlitzDoc/ThreadScheduler.htm>

<http://facultyweb.cs.wvu.edu/~phil/classes/blitz/BlitzDoc/ThreadScheduler.pdf>

2 Download the Files

Start by creating a new branch in your gitlab project named "a2". Checkout the new branch and then create a directory named "a2". Then download all the files from:

<http://facultyweb.cs.wvu.edu/~phil/classes/f19/509/a2>

into your directory. You should get the following files:

```
makefile
DISK
System.h
System.k
Runtime.s
Switch.s
List.h
List.k
Thread.h
Thread.k
Main.h
Main.k
Synch.h
Synch.k
Game.h
Game.k
Barber.h
Barber.k
SleepingBarberProblem.pdf
```

In this assignment, you will modify and hand in the following files:

```
Main.k
Synch.h
Synch.k
Game.h
Game.k
Barber.h
Barber.k
```

After getting the files, you should be able to compile all the code (as is) with the UNIX make command. The program executable we are building will be called "os". You can execute the first program by typing:

```
% make
% blitz -g os
```

It also builds "game" and "barber" that can be executed in a similar way. In the course of your experimentation, you may modify other files besides Synch.h, Synch.k, Main.k, Game.* and Barber.*. but the code you are required to write and turn in doesn't require any changes to the other files other than Synch.h, Synch.k, Main.k, Game.* and Barber.*. For example, you may wish to uncomment some of the print statements, to see what happens. *However, your final versions of the files you change must work with the other files, exactly as they are distributed.*

Be sure you copy all files. Even though there are similarities with some of the files used for assignment 1, there may be some subtle but important differences.

3 Study the Existing Code

The code provided in this assignment provides the ability to create and run multiple threads, and to control concurrency through several synchronization methods.

Start by looking over the System package. Focus on the material toward the beginning of file System.k, namely the functions:

```
print
printInt
printHex
printChar
printBool
nl
MemoryEqual
StrEqual
StrCopy
StrCmp
```

```
Min
Max
printIntVar
printHexVar
printBoolVar
printCharVar
printPtr
```

Get familiar with the printing functions; you'll be calling them a lot. Some are implemented in assembly code and some are implemented in KPL in the System package.

The following functions are used to implement the heap in KPL:

```
KPLSystemInitialize
KPLMemoryAlloc
KPLMemoryFree
```

Objects can be allocated on the heap and freed with the alloc and free statements. The HEAP implementation is very rudimentary in this implementation. In your kernel, you may allocate objects during start-up but after that, **YOU MUST NOT ALLOCATE OBJECTS ON THE HEAP!** Why? Because the heap might fill up and then what is a kernel supposed to do? Crash.

In this assignment, you should not allocate anything on the heap.

The following functions can be ignored since they concern aspects of the KPL language that we will not be using:

```
KPLUncaughtThrow
UncaughtThrowError
KPLIsKindOf
KPLSystemError
```

The Runtime.s file contains a number of routines coded in assembly language. It contains the program entry point and the interrupt vector in low memory. Take a look at it. Follow what happens when program execution begins at location 0x00000000 (the label "_entry"). The code labeled "_mainEntry" is included in code the compiler produces. The "_mainEntry" code will call the main function, which appears in the file Main.k.

In Runtime.s, follow what happens when a timer interrupt occurs. It makes an "up-call" to a routine called _P_Thread.TimerInterruptHandler. This name refers to "a function called TimerInterruptHandler in a package called Thread." (_P_Thread.TimerInterruptHandler is the name the compiler will give this function.)

All the code in this assignment assumes that no other interrupt types (such as a DiskInterrupt) occur. In Runtime.s, follow what would happen if another sort of interrupt should ever occur.

The KPL language will check for lots of error conditions, such as the use of a null pointer. Try changing the program to make this error. Follow in Runtime.s what happens when this occurs.

Next take a look at the List package. Read the header file carefully. This package provides code that implements a linked list. We'll use linked lists in this assignment. For example, the threads that are ready to run (and waiting for time on the CPU) will be kept in a linked list called the "ready list". Threads that become BLOCKED will sit on other linked lists. Also look over the List.k code file.

The key class in this assignment is named Thread, and it is located in the Thread package along with other stuff (see Thread.h, Thread.k). For each thread, there will be a single Thread object. Thread is a subclass of Listable, which means that each Thread object contains a next pointer and can be added to a linked list.

The Thread package is central and you should study this code thoroughly. This package contains one class (called Thread) and several functions related to thread scheduling and time-slicing:

```
InitializeScheduler ()
IdleFunction (arg: int)
Run (nextThread: ptr to Thread)
PrintReadyList ()
ThreadStartMain ()
ThreadFinish ()
FatalError (errorMessage: ptr to array of char)
SetInterruptsTo (newStatus: int) returns int
TimerInterruptHandler ()
```

FatalError is the simplest function. We will call FatalError whenever we wish to print an error message and abort the program. Typically, we'll call FatalError after making some check and finding that things are not as expected. FatalError will print the name of the thread invoking it, print the message, and then shut down. It will throw us into the BLITZ emulator command line mode. Normally, the next thing to do might be to type the "st" command (short for "stack"), to see which functions and methods were active.

(Of course the information printed out by the emulator will pertain to only the thread that invoked FatalError. The emulator does not know about threads, and it is pretty much impossible to extract information about other threads by examining bytes in memory.)

The next function to look at is SetInterruptsTo, which is used to change the "I" interrupt bit in the CPU. We can use it to disable interrupts with code like this:

```
... = SetInterruptsTo (DISABLED)
```

and we can use it to enable interrupts:

```
... = SetInterruptsTo (ENABLED)
```

This function returns the previous status. This is very useful because we often want to DISABLE interrupts (regardless of what they were before) and then later we want to return the interrupt status to whatever it was before. In our kernel, we'll often see code like:

```

var oldIntStat: int
...
oldIntStat = SetInterruptsTo (DISABLED)
...
oldIntStat = SetInterruptsTo (oldIntStat)

```

Next take a look at the Thread class. Here are the fields of Thread:

```

name: ptr to array of char
status: int
systemStack: array [SYSTEM\_STACK\_SIZE] of int
regs: array [13] of int
stackTop: ptr to void
initialFunction: ptr to function (int)
initialArgument: int

```

Here are the operations (i.e., methods) you can do on a Thread:

```

Init (n: ptr to array of char)
Fork (fun: ptr to function (int), arg: int)
Yield ()
Sleep ()
CheckOverflow ()
Print ()

```

Each thread is in one of the following states: JUST_CREATED, READY, RUNNING, BLOCKED, and UNUSED, and this is given in the status field. (The UNUSED status is given to a Thread after it has terminated. We'll need this in later assignments.)

Each thread has a name. To create a thread, you'll need a Thread variable. First, use Init to initialize it, providing a name.

Each thread needs its own stack and space for this stack is placed directly in the Thread object in the field called systemStack. Currently, this is an array of 1000 words, which should be enough. (It is conceivable our code could overflow this limit; there is a check to make sure we don't overflow this limited area.)

All threads in this assignment will run in System mode. Therefore the stack is called the "system stack". In later assignments, we'll see that this stack is used only for kernel routines. User programs will have their own stacks in their virtual address spaces, but we are getting ahead of ourselves.

The Thread object also has space to store the state of the CPU, namely the registers. Whenever a thread switch occurs, the registers will be saved in the Thread object. These fields (regs and stackTop) are used by the assembly code routine named Switch.

The Thread object also has space to store a pointer to a function (the initialFunction field) and an argument for this function (the initialArgument field). This pointer will point to the "main" function of this thread; this is the function that will get executed when this thread

begins execution. We are storing a pointer to the function because this is a variable: different threads may execute different functions.

We are also able to supply an initial argument to this thread, through the `initialArgument` field. This argument must be an integer. Often there will be several threads executing the same "main" function. The argument is a handy way to let each thread know what its role should be. For example, we might create 10 threads each using the same "main" function, but passing each thread a different integer (say, between 1 and 10) to let it know which thread it is.

After initializing a new Thread, we can start it running with the `Fork` method. This doesn't immediately begin the thread execution; instead it makes the thread `READY` to run and places it on the `readyList`. The `readyList` is a linked list of Threads. All Threads on the `readyList` have status `READY`. `ReadyList` is a global variable. There is another global variable named `currentThread`, which points to the currently executing Thread object; i.e., the Thread whose status is `RUNNING`.

The `Yield` method should only be invoked on the currently running thread. It will cause a switch to some other thread.

Follow the code in `Yield` closely to see what happens when a thread switch occurs. First, interrupts are disabled; we don't want any interference during a thread switch. The `readyList` and `currentThread` are shared variables and, while switching threads, we want to be able to access and update them safely. Then `Yield` will find the next thread from the `readyList`. (If there is no other thread, then `Yield` is effectively a `nop`.) Then `Yield` will make the currently running process `READY` (i.e., no longer `RUNNING`) and it will add the current thread to the tail end of the `readyList`. Finally, it will call the `Run` function to do the thread switch.

The `Run` method will check for stack overflow on the current thread. It will then call `Switch` to do the actual Switch.

`Switch` may be the most fascinating function you ever encounter! It is located in the assembly code file `Switch.s`, which you should look at carefully. `Switch` does not return to the function that called it. Instead, it switches to another thread. Then it returns. Therefore, the return happens to another function in another thread!

The only place `Switch` is called is from the `Run` function, so `Switch` returns to some invocation of the `Run` function in some other thread. That copy (i.e., invocation) of `Run` will then return to whoever called it. This could have been some other call to `Yield`, so we'll return to another `Yield` which will return to whoever called it.

And this is exactly the desired functionality of `Yield`! A call to `Yield` should give up the processor for a while, and eventually return after other threads have had a chance to execute.

`Run` is also called from `Sleep`, so we might be returning from a call to `Sleep` after a thread switch.

Question 1: How is everything set up when a thread is first created? How can we "return to a function" when we have not ever called it? Take a look at function `ThreadStartMain` in file `Thread.k` and look at function `ThreadStartUp` in file `Switch.s`.

Question 2: What happens when a thread is terminated? Take a look at `ThreadFinish` in file `Thread.k`. How is the thread put to sleep with no hope of ever being awakened? (No wonder they call it "Thread Death!")

Next, take a look at what happens when a Timer interrupt occurs while some thread is executing. This is an interrupt, so the CPU begins by interrupting the current routine's exe-

cution and pushing some state onto its (system) stack. Then it disables interrupts and jumps to the assembly code routine called `TimerInterruptHandler` in `Runtime.s`, which just calls the `TimerInterruptHandler` function in `Thread.k`.

In `TimerInterruptHandler`, we call `Yield`, which then switches to another thread. Later, we'll come back here, when this thread gets another chance to run. Then, we'll return to the assembly language routine which will execute a "ret" instruction. This will restore the state to exactly what it was before and the interrupted routine (whatever it was) will get to continue.

Note that this code maintains a variable called `currentInterruptStatus`. This is because it is rather difficult to query the "I" bit of the CPU. It is easier to just change the variable whenever a change to the interrupt status changes. We see this occurring in the `TimerInterruptHandler` function. Clearly interrupts will be disabled immediately after the interrupt occurs. And the `Yield` function will preserve the interrupt status. So when we return from `Yield`, interrupts will once again be disabled. Before returning to the interrupted thread, we set the `currentInterruptStatus` to `ENABLED`. (They must have been enabled before the interrupt occurred-or else it could not have occurred-so after we execute the "ret" instruction, the status will revert to what it was before, namely `ENABLED`.)

Now you are ready to start playing with and modifying the code! Please experiment with the code we have just discussed, as necessary to understand it.

4 Run the "SimpleThreadExample" Code

Execute and trace through the output of `SimpleThreadExample` in file `Main.k`.

In `TimerInterruptHandler` there is a statement

```
printChar ('\_')
```

which is commented out. Try uncommenting it. Make sure you understand the output.

Question 3: In `TimerInterruptHandler`, there is a call to `Yield`. Why is this there? Try commenting this statement? What happens. Make sure you understand how `Yield` works here.

5 Run the "MoreThreadExamples" Code

Trace through the output. Try changing this code to see what happens.

6 Implement the "Mutex" Class

In this part, you must implement the class `Mutex`. The class specification for `Mutex` is given to you in `Synch.h`:

```
class Mutex
  superclass Object
  methods
    Init ()
```

```

    Lock ()
    Unlock ()
    IsHeldByCurrentThread () returns bool
endClass

```

You will need to provide code for each of these methods. In `Synch.k` you'll see a behavior construct for `Mutex`. There are methods for `Init`, `Lock`, `Unlock`, and `IsHeldByCurrentThread`, but these have dummy bodies. You'll need to write the code for these four methods.

You will also need to add a couple of fields to the class specification of `Mutex` to implement the desired functionality.

How can you implement the `Mutex` class? Take a close look at the `Semaphore` class; your implementation of `Mutex` will be quite similar.

First consider the `IsHeldByCurrentThread` method, which may be invoked by any thread. The code of this method will need to know which thread is holding a lock on the mutex; then it can compare that to the `currentThread` to see if they are the same. So, you might consider adding a field (perhaps called `heldBy`) to the `Mutex` class, which will be a pointer to the thread holding the mutex. Of course, you'll need to set it to the current thread whenever the mutex is locked. You might use a null value in this field to indicate that no thread is holding a lock on the mutex.

When a lock is requested on the mutex, you'll need to see if any thread already has a lock on this mutex. If so, you'll need to put the current process to sleep. For putting a thread to sleep, take a look at the method `Semaphore.Down`. At any one time, there may be zero, one, or many threads waiting to acquire a lock on the mutex; you'll need to keep a list of these threads so that when an `Unlock` is executed, you can wake up one of them. As in the case of `Semaphores`, you should use a FIFO queue, waking up the thread that has been waiting longest.

When a mutex lock is released (in the `Unlock` method), you'll need to see if there are any threads waiting to acquire a lock on the mutex. You can choose one and move it back onto the `readyList`. Now the waiting thread will begin running when it gets a turn. The code in `Semaphore.Up` does something similar.

It is also a good idea to add an error check in the `Lock` method to make sure that the current thread asking to lock the mutex doesn't already hold a lock on the mutex. If it does, you can simply invoke `FatalError`. (This would probably indicate a logic error in the code using the mutex. It would lead to a deadlock, with a thread frozen forever, waiting for itself to release the lock.) Likewise, you should also add a check in `Unlock` to make sure the current thread really does hold the lock and call `FatalError` if not. You'll be using your `Mutex` class later, so these checks will help your debugging in later assignments.

The function `TestMutex` in `Main.k` is provided to exercise your implementation of `Mutex`. It creates 7 threads that compete vigorously for a single mutex lock.

7 Implement the Producer-Consumer Solution

Your textbook, Silberschatz, et al., contains a discussion of the Producer-Consumer problem, including a solution. Implement this in KPL using the classes `Mutex` and `Semaphore`. Deal with multiple producers and multiple consumers, all sharing a single bounded buffer.

The Main package contains some code that will serve as a framework. The buffer is called `buffer` and contains up to `BUFFER_SIZE` (e.g., 5) characters. There are 5 producer processes, each modeled by a thread, and 3 consumer processes, each modeled by a thread. Thus, there are 8 threads in addition to the main thread that creates the others.

Each producer will loop, adding 5 characters to the buffer. The first producer will add five 'A' characters, the second producer will add five 'B's, etc. However, since the execution of these threads will be interleaved, the characters will be added in a somewhat random order.

8 Implement the Dining Philosopher's Solution Using a Monitor

A starting framework for your solution is provided in `Main.k`. Each philosopher is modeled with a thread and the code we've provided sets up these threads. The synchronization will be controlled by a "monitor" called `ForkMonitor`.

The code for each thread/philosopher is provided for you. Look over the `PhilosophizeAndEat` method; you should not need to change this code.

The monitor to control synchronization between the threads is implemented with a class called `ForkMonitor`. The following class specification of `ForkMonitor` is provided:

```
class ForkMonitor
  superclass Object
  fields
    status: array [5] of int      -- For each philosopher: HUNGRY,
                                -- EATING, or THINKING

  methods
    Init ()
    PickupForks (p: int)
    PutDownForks (p: int)
    PrintAllStatus ()
endClass
```

You'll need to provide the code for the `Init`, `PickupForks` and `PutDownForks` methods. You'll also need to add additional fields and perhaps even add another method.

The code for `PrintAllStatus` is provided. You should call this method whenever you change the status of any philosopher. This method will print a line of output, so you can see what is happening.

How can you proceed? You'll need a mutex to protect the monitor itself. There are two main methods (`PickupForks` and `PutDownForks`) which are called by the philosopher threads. Upon beginning each of these methods, the first thing is to lock the monitor mutex. This will ensure that only one thread at a time is executing within the monitor. Just before each of these methods returns, it must unlock the monitor (by unlocking the monitor's mutex) so that other threads can enter the monitor code.

You'll also need to use the Condition class, which is provided in the Synch package. (The Condition class uses the class Mutex, so it is assumed that you've finished and tested the Mutex class.)

The BLITZ emulator has a number of parameters and one of these is how often a timer interrupt occurs. The default value is every 5000 instructions. You might try changing this parameter to see how it affects your programs behavior. To change the simulation parameters, type the sim command into the emulator. This command will give you the option to create a file called

```
.blitzrc
```

After creating this file, you can edit it by hand. The next time you run the emulator, it will use this new value. Also note that too small a value-like 1000-will cause the program to hang.

Question 4: What do you suppose causes this effect?

Question 5: Find and read about "Hoare semantics" and "Mesa-style semantics". The comments in the code in Synch.k say that this version implements "Mesa-style" semantics. Is this the same or different from Hoare semantics? If different, what is the difference? Quote the code that implements the "Mesa-style semantics."

9 An Example of Correct Output

The following files contain an example of what correct output should look like:

```
DesiredOutput1.pdf
DesiredOutput2.pdf
DesiredOutput3.pdf
DesiredOutputGame.pdf (for the next two steps)
DesiredOutputBarber.pdf
```

10 Implement the Gaming Parlor Problem (50 points)

Here is the problem scenario: groups of customers come in to a "gaming parlor" to play games. They go to the front desk to obtain one or more dice, which are used by the group while they are playing their game, and then returned to the front desk. The front desk is in charge of lending out the dice and collecting them after each game is finished.

The gaming parlor owns only 8 dice, which are available at the front desk before the first group comes in.

The customers can play the following games. Listed after each game in parentheses is the number of dice required to play that game.

```
Backgammon (4)
Risk (5)
Monopoly (2)
Pictionary (1)
```

You should model the front desk as a monitor with the following entry methods:

```
Request (numberOfDice: int)
Return (numberOfDice: int)
```

Model each group of customers as a thread. When a group is ready to play, it must obtain the necessary number of dice. If the required number of dice is not available, then the group (i.e., the thread) must wait. You might use a condition variable that "more dice have become available."

You should model the following eight different groups. Each group plays one game, as shown below, but each group plays its game 5 times. Each group must return their dice after each game and then re-acquire the dice before playing again.

```
A - Backgammon
B - Backgammon
C - Risk
D - Risk
E - Monopoly
F - Monopoly
G - Pictionary
H - Pictionary
```

To simulate playing the game, simply call the Yield method.

This problem is a generalization of the problem of resource allocation where (1) there are a number of resources (dice) but each is identical; (2) every requesting process needs a one or more units of the resource, (3) each requesting thread knows how many units it will need before requesting any units and that info is included in the request, (4) all units are returned before any further requests are made.

In a monitor, each method is either an "entry" method or a local method. A monitor will always have one Mutex lock, which is associated with entering the monitor. Every monitor will also have zero or more condition variables, as required by the particular application.

An entry method must always begin by locking the monitor's mutex. Before it returns, it must always unlock the mutex. Within the entry method, the code may Signal some of the condition variables and may Wait on condition variables. (In some applications, it may make sense to Broadcast to a condition variable.) However, once inside of an entry method, the code should never touch the monitor's mutex lock, or try to look inside the condition variables.

You should model the front desk as a monitor and you should use condition variables instead of semaphores.

If deadlock occurs, the program will freeze up and some requests for dice will go unsatisfied forever. This is a disaster! Regardless of the order in which the groups make their requests, your solution should be structured such that deadlock can never occur.

Your solution must not be subject to any race conditions. In other words, regardless of the order in which the groups make their requests and return their dice, each die must never be allocated to more than one group at a time. It should never be the case that groups are allowed

to proceed when there are too few dice. Likewise, if a group has returned its dice, other groups which are waiting must be allowed to proceed once enough dice have become available.

Starvation can occur if it is possible that one thread can be delayed infinitely by other threads' requests. If the game parlor problem is extended to assume that each group will continue to play game after game forever, then starvation might be possible, if you are not careful in your solution. If some group X comes in requesting a lot of dice, it will be made to wait until enough dice are available. If it is possible that other groups can come in, request a small number of dice, and have their requests granted, then group X might get delayed forever, since there are never enough dice at the front desk at once to satisfy group X's needs. In other words, it might be possible for starvation of X to occur. In your solution starvation should not be possible. (Of course, with each group limited to playing only 5 games, all the outstanding dice will always get returned eventually and starvation can never occur, but your solution should also avoid starvation in the presence of a never-ending sequence of Request and Return operations.)

To verify that your code is working, please insert print statements to produce output like this...

```
Initializing Thread Scheduler...
A requests 4
-----Number of dice now avail = 8
A proceeds with 4
-----Number of dice now avail = 4
B requests 4
-----Number of dice now avail = 4
B proceeds with 4
-----Number of dice now avail = 0
D requests 5
-----Number of dice now avail = 0
E requests 2
-----Number of dice now avail = 0
A releases and adds back 4
-----Number of dice now avail = 4
B releases and adds back 4
-----Number of dice now avail = 8
C requests 5
-----Number of dice now avail = 8
H requests 1
-----Number of dice now avail = 8
B requests 4
-----Number of dice now avail = 8
D proceeds with 5
-----Number of dice now avail = 3
...etc...
```

This output makes it fairly easy to see what the program is doing and verify that it is

correct.

Below is a method you should use to produce your output. Feel free to copy this method straight into your program.

```
behavior GameParlor
```

```
... Other methods...
```

```
method Print (str: String, count: int)
--
-- This method prints the current thread's name and the arguments.
-- It also prints the current number of dice available.
--
  print (currentThread.name)
  print (" ")
  print (str)
  print (" ")
  printInt (count)
  nl ()
  print ("-----Number of dice now avail = ")
  printInt (numberDiceAvail)
  nl ()
endMethod
```

```
endBehavior
```

This method would be called in several places...

- At the beginning of the Request method: self.Print ("requests", numNeeded)
- At the end of the Request method: self.Print ("proceeds with", numNeeded)
- In the Return method: self.Print ("releases and adds back", numReturned)

11 Implement the Sleeping Barber Problem (50 points)

An older edition of the Tanenbaum textbook describes the "Sleeping Barber" problem and gives a solution in "C". (This material can also be found in a separate file called SleepingBarber-Problem.pdf in the a2 directory on the class web site.)

Translate this into a working KPL program. You'll also need to create some code to print out what happens when you run the program.

What will you do for the "cut_hair" and the "get_haircut" methods? You'll need at least one print statement in each routine, but be careful: they both should run at more-or-less the same time.

One barber is okay, but how many customers will you model? Will you just throw a bunch of customers at the barbershop all at once? This might not test your code very well.

How long will the haircut last? You can implement waiting with a busy loop, such as

```
for i = 1 to 100
  .. want to yield here?...
endFor
```

12 What to Hand In

Complete all the above steps you are required to complete. In a text or word processed document, answer all five questions asked in the above text.

All your files for this assignment must be in the “a2” directory of your “a2” branch. The following are the files you should have modified. All other files should be unmodified.

Synch.h, Synch.k, Main.k, Game.* and Barber.*

Capture the output for sections 6, 7, 8, 10 and 11 in separate files named **section-X** where X is the section number. (script(1) is a great program to use for this purpose.) Add and commit those files to your “a2” directory. Remember to push all your changes to gitlab.

For the Barbers implements, write up a convincing argument that your implementation is correct.

You should turn in the following in a printed document collection:

1. A cover sheet with your name, this class number and assignment number.
2. Answers to the five questions.
3. Your convincing argument
4. If you want comments on your code, you may print your code and include it with your turn answers. a2ps(1) is a great program to produce readable output.

There should be no hand written work in the documents submitted.

13 Basis for Grading

The code you submit will be graded on both correctness and style. Correctness means that the code must work and not contain bugs. Style means the the code is neat, well organized, well commented and easy to read. Indentation and formatting should match the distributed code.